

Degrees of Software Correctness

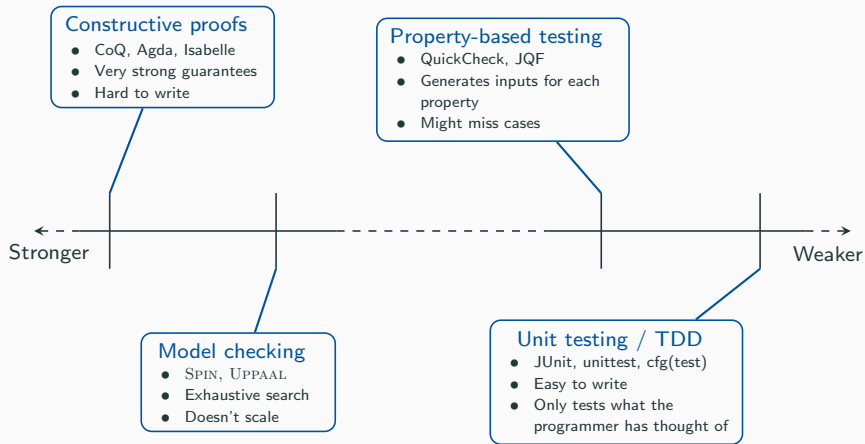
Uniting the Spectrum of Verification

Thomas Ekström Hansen

10th November 2023

- We generally want software to be correct
- Many approaches to this
 - Test Driven Development (TDD)
 - Mutation testing
 - Property-based testing
 - Model Checking
 - Type theory
 - Mathematical abstractions and proofs
- No one-size-fits-all
 - So compromises are made

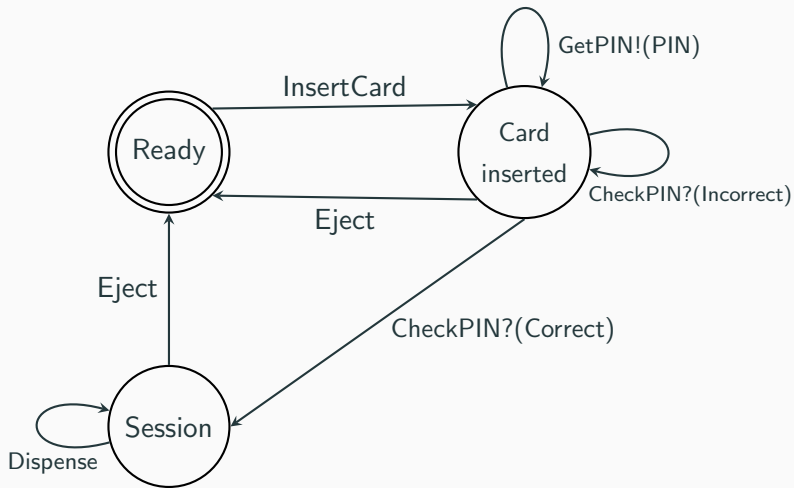
Spectrum of Verification



The eternal problem with verification systems

- All verification systems face the same problem: ergonomics
- If the system is not unobtrusive, people are unlikely to use it
- Our hope is that our approach tries to not get in the way by being part of the program
 - Compiler and type-system assist you rather than hinder
 - Verify as you go along, tuning the strictness as necessary
 - Escape hatches provided for prototyping

Visuals can really help



State machine of an ATM

Using types only gets you part of the way there

```
data State    = Ready | CardInserted | Session

data CheckPINRes = Incorrect | Correct

data CMD : (ty : Type) -> State -> (ty -> State) -> Type
  where
    InsertCard : CMD () Ready (const CardInserted)
    Dispense    : CMD () Session (const Session)
    GetPIN      : CMD () CardInserted (const CardInserted)

    CheckPIN    : (Vect 4 Nat)
                  -> CMD CheckPINRes CardInserted
                  (\res => case res of
                      Incorrect => CardInserted
                      Correct  => Session)

    GetAmount   : CMD Nat state (const state)

    EjectCard   : CMD () state (const Ready)
```

Using types only gets you part of the way there

Accepted by the compiler:

```
atmProg : CMD ()
    Ready (const Ready)
atmProg =
    do InsertCard
        pin <- GetPIN
        Correct <- CheckPIN pin
            | Incorrect => EjectCard
        amt <- GetAmount
        Dispense amt
        EjectCard
```

Rejected by the compiler:

```
badProg : CMD ()
    Ready (const Ready)
badProg =
    do InsertCard
        pin <- GetPIN
        Correct <- CheckPIN pin
            | Incorrect => InsertCard
        amt <- GetAmount
        Dispense amt
        -- forgot to eject
```

Disadvantages of the CMD-like approach


- Not the most ergonomic
 - It is a Domain Specific Language
 - So it needs an interpreter; it is not directly runnable code
 - The type declarations can get lengthy
- Composition is currently undefined
 - The theory it builds on does not currently support composition
 - But we often want to compose code and functionality

Benefits of a multifaceted approach

1. Adaptability — right tool for the job
2. Speed — can trade speed for level of verification
3. **Coherence** — all done in one system
 - No need to translate to model-checking tool
 - Specification lives alongside model lives alongside implementation
 - The implementation is just there; it *is* runnable code
 - Parts can be verified independently *while* combining into an overall system

Example: ARQ – Overview

- Automated Repeat reQuest (of course...)
- Sender sends one packet at a time
- Wait for the receiver to acknowledge they are ready for the next packet
- Resend if no acknowledgement is received within a given timeout



Sample ARQ
transmission

Example: ARQ – Verification challenges

Different parts necessitate different amounts of verification

- Data verification is relatively simple
 - Proofs might be feasible.
- Potentially infinite packet numbers
 - Proofs will be tedious
 - Not fit for model checking
- Network interface possibly outwith our control
 - Likely using a library or similar
 - Cannot control how interfaces behave
- Other half likely not even physically local
 - Interfacing with someone else's code and/or infrastructure
 - Impossible to reason about beyond handling responses

Example: ARQ – Our approach

- Compression can be proven to be correctly implemented
 - Types for compressed and uncompressed data
 - Link the two via ``compress`` and ``uncompress`` functions
 - The compiler can then help us implement those correctly
- QuickCheck can simulate network behaviour
 - Generate an arbitrary packet
 - Simulate sending with a chance of failure
 - Check that 10 or 100 different sequences all get delivered
 - All done at compile-time, using the same datatypes as the implementation
- Some unit tests could be added for the network implementation
 - Interaction with external libraries, so unit tests is best we can do
 - Confidence from QuickCheck increased by these passing

Example: ARQ – Benefits of our approach

- All in the same system
- We do not need to trust we have modelled the problem correctly, we can test that!
- Can use as stronger or weaker verification as wanted/required
- There is a link between the verified and implemented model: the datatypes!
- The compiled program is an executable version of our verified implementation!

- Many verification tools exist, none of them cover enough on their own
- Instead of “competing”, we combine the systems to work together
- Hopefully this will lead to wider adoption and better whole-system soundness

Questions?



github.com/CodingCellist/talks