Increasing Confidence in Types

Thomas Ekström Hansen
6th March 2024

Overview

- Many systems exhibit Finite State Machine (FSM)-like behaviour
  - These can be neatly modelled with Dependent Types
- ATM example of this from Edwin Brady's "Type Driven Development with Idris" book
  - Turns out this example is subtly incorrect
- How can we spot this? Can we increase confidence in our type-level modelling without having to "just get it right"?
- Spoilers: How does QuickCheck fit in with dependent types?
- No silver bullet, but we can hopefully catch errors faster and have some guarantees that our model behaves as intended

The ATM state machine
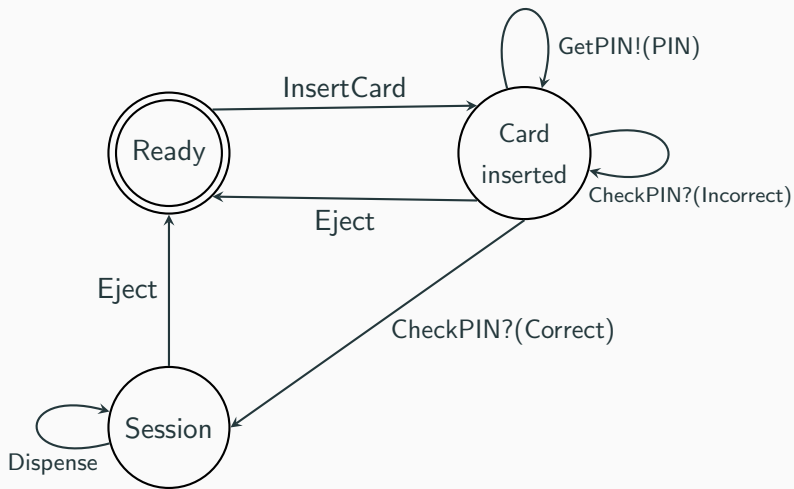
State machine of an ATM

Indexed State Monads (ISMs)

```
data ATMSt =  Ready | CardInserted | Session

data CheckPINRes  =  Incorrect | Correct

data ATMOp : (ty : Type) -> ATMSt -> (ty -> ATMSt) -> Type
  where
  Insert : ATMOp () Ready (const CardInserted)

  CheckPIN : (pin : Nat)
          -> ATMOp CheckPINRes CardInserted
                     (\case Incorrect => CardInserted
                            Correct => Session)

  GetAmount  : ATMOp Nat state (const state)
  Dispense :  (amt : Nat)
          -> ATMOp () Session (const Session)
  Eject  : ATMOp () state (const Ready)
```

Using the operations

- Using the ISM operations requires another ISM, defining `pure`, `op`, `bind`, and `seq`

```
data ATM : (t : Type) -> ATMSt -> (t -> ATMSt) -> Type
  where
  Pure : (x : t) -> ATM t (stFn x) stFn
  Op : ATMOp t st st' -> ATM t st st'
  (>>=) :  ATM a s1 s2f -> ((x : a) -> ATM b (s2f x) s3f)
        -> ATM b s1 s3f
  (>>) :  ATM () s1 s2f -> (ATM b (s2f ()) s3f)
        -> ATM b s1 s3f
```

Programming with ISMs

- We declare our intended start and end state in the type
- And the type-checker verifies that we don't use commands incorrectly

```
prog : ATM () Ready (const Ready)
prog = do
  Op Insert
  Correct <- Op $ CheckPIN 1234
    | Incorrect => <...>
  amount <- Op GetAmount
  Op $ Dispense amount
  Op Eject
```

Using types only gets you part of the way there

Rejected by the type-checker:

```
badProg : ATM ()
               Ready (const Ready)
badProg = do
  Op Insert
  let pin = 1234
  Correct <- Op $ CheckPIN pin
    | Incorrect => InsertCard
  amt <- Op GetAmount
  Op $ Dispense amt
  -- We never Eject, so we
  -- never come back to
  -- `Ready'
```

Accepted by the type-checker:

```
loopProg : CMD ()
             Ready (const Ready)
loopProg = do
    Op InsertCard
    let pin = 4321
    loopIncorrect pin
  where
    loopIncorrect : Nat -> ATM ()
                      CardInserted
                      (const Ready)
    loopIncorrect p = do
      Incorrect <- Op $ CheckPIN p
        | Correct => -- <...>
      loopIncorrect p
```
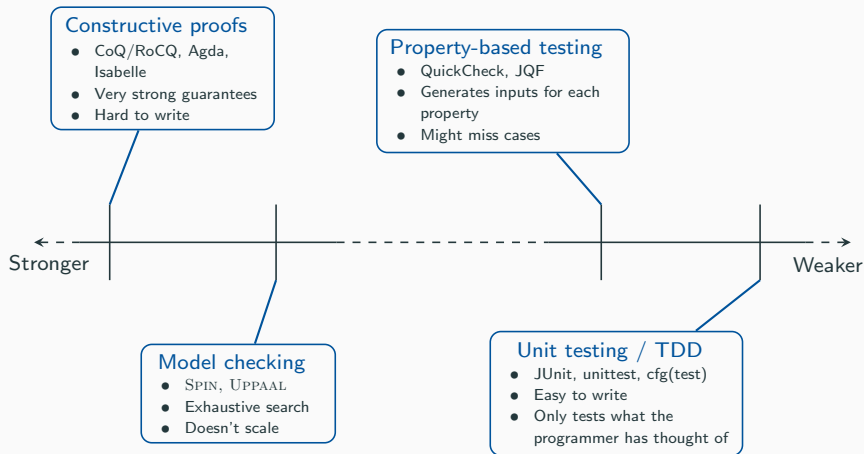
Why is this a problem?

- We expect an ATM to reject the card after 3 PIN attempts
  - Not to be permanently unavailable if we retry forever
- However, the programmer is unlikely to catch this
- The model looks correct and rigorous, after all
- Programming with it will catch most errors
- And the type-checker is happy with our sequence of operations

How do we solve this?

- We could spot the issue when it happens
  - Someone will (hopefully) spot the issue during development
  - Or, worst case, spot it when it happens after deployment
- And then we update our model and everything is good
- Why not try to spot it *automatically* before either of those?
- Modelling can clearly go wrong, so how do we increase our confidence in the models? Who type-checks the types?

Constructive proofs
- CoQ/RoCQ, Agda, Isabelle
- Very strong guarantees
- Hard to write

Property-based testing
- QuickCheck, JQF
- Generates inputs for each property
- Might miss cases

Stronger — — — — — — — — — — — — — — — — — — — Weaker

Model checking
- SPIN, UPPAAL
- Exhaustive search
- Doesn't scale

Unit testing / TDD
- JUnit, unittest, cfg(test)
- Easy to write
- Only tests what the programmer has thought of

The eternal problem with verification systems

- All verification systems face the same problem: ergonomics
- If the system is obstructive, or even just perceived as such, people are unlikely to use it
  - This is especially true for complex systems
  - "Fighting with the Rust borrow-checker"
  - "I'm experienced enough to write safe C/C++"
- Where does Idris fit in here?
  - Allows us to target all areas of the spectrum
  - Compiler and type-system assist you rather than hinder
  - Verify as you go along, tuning the strictness as necessary
  - Unit tests are not thorough enough, so QuickCheck seems like a good middle ground
  - Dependent types allow us to run the tests at compile time, and quantities to erase their results at runtime!

How do you generate a dependent type?

- QuickCheck's bread and butter is `Arbitrary`
  - Define how to generate an instance of a type, given some pseudorandom number generator state
  - Reasonably straightforward for random numbers, picking an element, and structures where the type of the constructors are known at generation-time
- However, our types are *dependent*
- So we cannot know the exact type at generation time
  - We can know *a* type, but not all. For example, `Vect 3 Nat` is trivial: `[!arbitrary, !arbitrary, !arbitrary]`
  - The problem is `Vect n Nat`
  - Or even `Arbitrary t => Vect n t`

- The solution is more dependent types!
- Specifically: dependent pairs

```
record DPair a (p : a -> Type) where
  constructor MkDPair
  fst : a
  snd : p fst
```

- As long as we know how to generate an `Arbitrary a`, we can generate an `Arbitrary (x : a ** p x)`
  - (The ** syntax is sugar for DPair / MkDPair depending on the context)

Plumbing for operations

```idris
record OpRes (resT : Type) (currSt : ATMState)
             (nsFn : resT -> ATMState) where
  constructor MkOpRes

  -- The operation
  op : ATMOp resT currSt nsFn

  -- The result of the operation
  res : resT

  -- Results must be `Show`-able for QC to work
  rShow : Show resT
```

Tracing ATMs

```
record TraceStep where
  constructor MkTS

  -- The `ATMOp`, along with some result,
  -- which took us to the traced state
  opRes : OpRes rT aSt aStFn

  -- The `ATMState` we ended up in
  resSt : ATMState

-- A bounded sequence of trace steps
data ATMTrace : ATMState -> Nat -> Type where
  MkATMTrace :  (initSt : ATMState)
             -> {bound : Nat}
             -> (trace : Vect bound TraceStep)
             -> ATMTrace initSt bound
```

Generating arbitrary OpRes

```
{currSt : ATMState} ->
Arbitrary (resT : _ ** nsFn : resT -> ATMState ** OpRes resT currSt nsFn)
where
  arbitrary {currSt=Ready} =
    pure (_ ** _ ** MkOpRes Insert () %search)

  arbitrary {currSt=CardInserted} = do
    -- we need _a_ PIN, even though we control the result
    let arbPIN = 0
    let op1 = (_ ** _ ** MkOpRes (CheckPIN arbPIN) Correct %search)
    let op2 = (_ ** _ ** MkOpRes (CheckPIN arbPIN) Incorrect %search)
    let op3 = (_ ** _ ** MkOpRes Eject () %search)
    -- can adjust the frequencies of getting the PIN wrong
    frequency $ [(1, pure op1), (4, pure op2), (1, pure op3)]

  arbitrary {currSt=Session} = do
    arbAmount <- arbitrary
    let op1 = (_ ** _ ** MkOpRes (Dispense arbAmount) () %search)
    let op2 = (_ ** _ ** MkOpRes Eject () %search)
    oneof $ map pure [op1, op2]
```

15

Properties of the ATM

- As you can see, generating arbitrary ATM steps is a bit more involved than non-dependent types, but it is doable

- And now that we have that, we can specify properties like "Within 5 state-transitions, we should be back in `Ready`"
  ```
  public export 0
  PROP_eventuallyReady : Fn (ATMTrace Ready 5) Bool
  PROP_eventuallyReady =
    MkFn (\case (MkATMTrace _ trace) =>
               elem Ready (map (.resSt) trace))
  ```

- *And* we can test it at compile-time
  ```
  public export 0
  EventuallyReady_OK : So (QuickCheck
                          PROP_eventuallyReady)
  EventuallyReady_OK = Oh
  ```

Model, verification, and implementation all-in-one

- With most verification tools, we have to translate between
  models
    - Spec, model, and implementation are independent
- This facilitates translation mistakes
    - Might think we're verifying the same thing, when in actual fact
      the semantics have changed between representations
- In our case, the specification *is* the model; *everywhere*

```
trace :  (steps : Nat) -> (currSt : ATMState)
      -> Gen (Vect steps TraceStep)
trace 0 _ = pure []
trace (S k) currSt =
  the (Gen (x ** y ** OpRes x currSt y)) arbitrary
  >>=
  \case (_ ** nsFn ** opR@(MkOpRes _ res _)) =>
     do let nextState = nsFn res  -- ! nsFn from ISM
        pure $ (MkTS opR nextState) :: !(trace k nextState)
```

QuickCheck spots the error!

- If we try to type-check the file we get:

```
-- Error: While processing right hand side of
--        EventuallyReady_OK. When unifying:
   So True
-- and:
   So (QuickCheck PROP_eventuallyReady)
-- Mismatch between: True and False.
```

- And if we investigate by running QC at the REPL, the error is exactly the fault in the model:

```
MkQCRes (Just False) <log> """
Falsifiable, after 4 tests:
Starting @ Ready:
[ (<ATMOp 'Insert ~ ()'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
]\n"""
```

Fixing things

- Now that we know there's an error, we can fix things!

```
data ATMState
  = Ready
  | CardInserted Nat  -- track #retries
  | Session
<...>
CheckPIN : (pin : Nat)
  -> ATMOp PINok (CardInserted (S tries))
       (\case Correct => Session
              Incorrect => ifThenElse (isZero tries)
                             Ready
                             (CardInserted tries))
```

- Carrying this through to the generators, our QC passes: file reloads successfully, the REPL reports

```
> QuickCheck PROP_eventuallyReady
MkQCRes (Just True) <log> "OK, passed 100 tests"
```

Benefits of a multifaceted approach

1. Adaptability — being able to use different tools
2. Speed — can trade speed for level of verification
   - This isn't about proving things, it is about increasing confidence in our typed models
3. **Coherence** — all done in one system
   - No need to translate to model-checking tool
   - Specification lives alongside model lives alongside implementation
   - The implementation is just there; it *is* runnable code
   - Parts can be verified independently *while* combining into an overall system

Further work

- As neat as this is, it is still convoluted to write the types, generators, etc.
- There are an abundance of FSM-like systems — the ARQ protocol, pick-and-place machines — which we plan to model
- This should hopefully reveal common patterns, which we can then factor out and automate large parts of this

Questions?

Slides

github.com/CodingCellist/talks