

Increasing Confidence in Types

Thomas Ekström Hansen

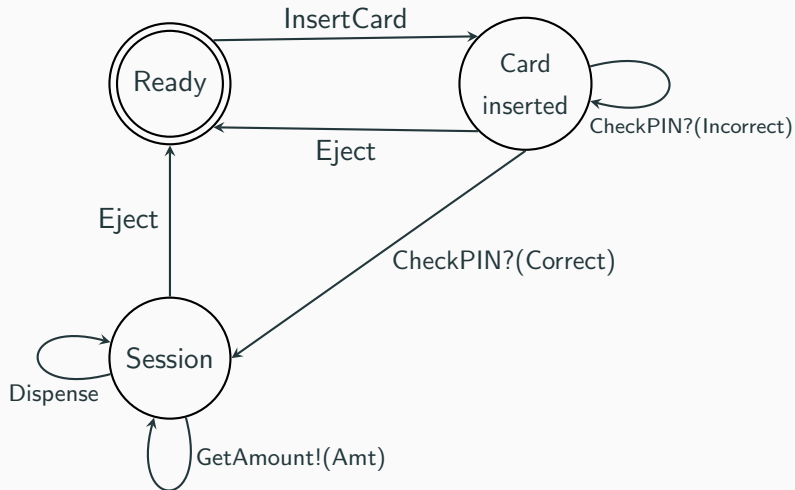
6th March 2024

- Many systems exhibit Finite State Machine (FSM)-like behaviour
- These can be modelled using dependent types
 - How do we know the model's behaviour is correct?
- How does QuickCheck fit in with dependent types?

This is not a proof technique

But hopefully, it helps us catch errors faster and provides guarantees that our model behaves as intended

The ATM state machine



State machine of an ATM

Indexed State Monads (ISMs)

```
data ATMSt = Ready | CardInserted | Session

data CheckPINRes = Incorrect | Correct

data ATMOp : (ty : Type) -> ATMSt -> (ty -> ATMSt) -> Type
  where
    Insert : ATMOp () Ready (const CardInserted)

    CheckPIN : (pin : Nat)
      -> ATMOp CheckPINRes CardInserted
      (\case Incorrect => CardInserted
        Correct => Session)

    GetAmount : ATMOp Nat Session (const Session)
    Dispense : (amt : Nat)
      -> ATMOp () Session (const Session)
    Eject : ATMOp () state (const Ready)
```

Using the operations

- Using the ISM operations requires another ISM, defining pure, op, bind, and seq

```
data ATM : (t : Type) -> ATMSt -> (t -> ATMSt) -> Type
where
  Pure : (x : t) -> ATM t (stFn x) stFn
  Op : ATMOp t st st' -> ATM t st st'
  (>>=) : ATM a s1 s2f -> ((x : a) -> ATM b (s2f x) s3f)
    -> ATM b s1 s3f
  (>>) : ATM () s1 s2f -> (ATM b (s2f ()) s3f)
    -> ATM b s1 s3f
```

Programming with ISMs

- We declare our intended start and end state in the type
`prog : ATM () Ready (const Ready)`
- And the type-checker verifies that we don't use commands incorrectly

```
prog = do                                     -- We start in Ready
  Op Insert ----- Ready to CardInserted
  Correct <- Op $ CheckPIN 1234 --- CI to Session
    | Incorrect => <...> ----- (or stay in CI)
  amount <- Op GetAmount ----- Stay in Session
  Op $ Dispense amount ----- Stay in Session
  Op Eject ----- Return to Ready
```

Using types only gets you part of the way there

Rejected by the type-checker:

```
badProg : ATM ()
          Ready (const Ready)

badProg = do
  Op Insert
  let pin = 1234
  Correct <- Op $ CheckPIN pin
    | Incorrect => InsertCard
  amt <- Op GetAmount
  Op $ Dispense amt
  -- We never Eject, so we
  -- never come back to
  -- `Ready'
```

Accepted by the type-checker:

```
loopProg : ATM ()
          Ready (const Ready)

loopProg = do
  Op InsertCard
  let pin = 4321
  loopIncorrect pin
where
  loopIncorrect : Nat -> ATM ()
                                CardInserted
                                (const Ready)

  loopIncorrect p = do
    Incorrect <- Op $ CheckPIN p
    | Correct => -- <...>
    loopIncorrect p
```

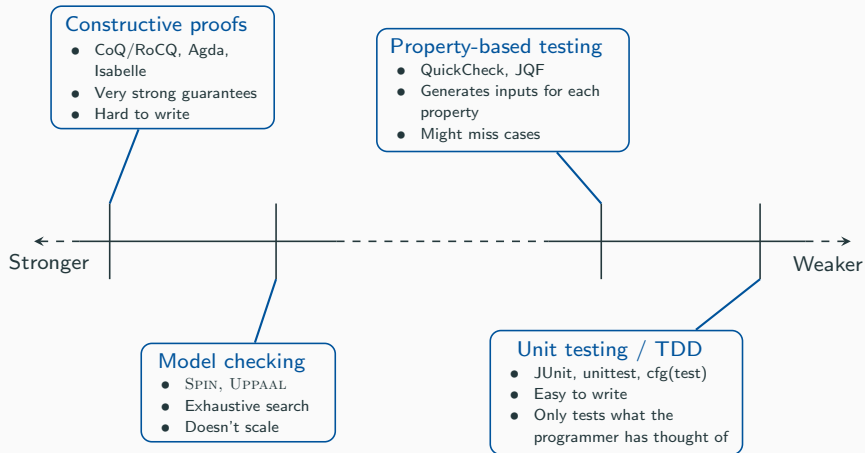

Why is this a problem?

- We expect an ATM to reject the card after 3 PIN attempts
 - Not to be permanently unavailable if we retry forever
- However, the programmer is unlikely to catch this
- The model looks correct and rigorous, after all
- Programming with it will catch most errors
- And the type-checker is happy with our sequence of operations

How do we solve this?

- We could spot the issue when it happens
 - Someone will (hopefully) spot the issue during development
 - Or, worst case, spot it when it happens after deployment
- And then we update our model and everything is good
- Why not try to spot it *automatically* before either of those?
- Modelling can clearly go wrong, so how do we increase our confidence in the models? Who type-checks the types?

Spectrum of Verification



The eternal problem with verification systems

- All verification systems face the same problem: ergonomics
- If the system is obstructive, or even just perceived as such, people are unlikely to use it
- This is especially true for complex systems
 - “Fighting with the Rust borrow-checker”
 - “I’m experienced enough to write safe C/C++”
 - “I’m experienced enough to get the types right”

Where does Idris fit in?

- General-purpose with dependent types, allowing us to program to different areas of the spectrum
- Compiler and type-system assist rather than hinder
- Verify as you go along, tuning the strictness as necessary
- Unit tests are not thorough enough, so QuickCheck seems like a good middle ground
- Dependent types allow us to run the tests at compile time, and quantities to erase their results at runtime!

How do you generate a type?

- QuickCheck's bread and butter is **Arbitrary**
- Define how to generate an instance of a type, given some pseudorandom number generator state
- We know how to do this for random numbers, picking an element, or structures where the type of the constructors are known at generation-time

How do you generate a dependent type?

- Our types can depend on values
- So we cannot know the exact type at generation time
- We can know *a* type, but not all. For example, `Vect 3 Nat` is trivial: `[!arbitrary, !arbitrary, !arbitrary]`
 - The problem is `Vect n Nat`
 - Or even `Vect n t`

Arbitrary dependent types

- The solution is more dependent types!
- Specifically: dependent pairs

```
record DPair a (p : a -> Type) where
  constructor MkDPair
  fst : a
  snd : p fst
```

- As long as we know how to generate an ``Arbitrary a``, we can generate an ``Arbitrary (x : a ** p x)``
 - (The `**` syntax is sugar for `DPair / MkDPair` depending on the context)

Arbitrary vectors

- Provided we know how to generate the elements, we generate *some* length

```
Arbitrary a => Arbitrary (n : Nat ** Vect n a) where
  arbitrary = do
    len <- arbitrary
```

- And then generate that many arbitrary elements

```
    vect <- nArbitrary len
    pure (len ** vect)
where
  nArbitrary : (n : Nat) -> Gen (Vect n a)
  nArbitrary 0 = []
  nArbitrary (S k) = !arbitrary :: nArbitrary k
```

Arbitrary ATMs?

- Can we do a similar thing for $ATMOp$ and ATM ?
- Yes, but we need some (dependent) plumbing first

Plumbing for operations

```
record OpRes (resT : Type) (currSt : ATMState)
    (nsFn : resT -> ATMState) where
  constructor MkOpRes

-- The operation
op : ATMOp resT currSt nsFn

-- The result of the operation
res : resT

-- Results must be `Show`-able for QC to work
rShow : Show resT
```

Tracing ATMs

```
record TraceStep where
  constructor MkTS

  -- The `ATMOp`, along with some result,
  -- which took us to the traced state
  opRes : OpRes rT aSt aStFn

  -- The `ATMState` we ended up in
  resSt : ATMState

  -- A bounded sequence of trace steps
data ATMTrace : ATMState -> Nat -> Type where
  MkATMTrace : (initSt : ATMState)
    -> {bound : Nat}
    -> (trace : Vect bound TraceStep)
    -> ATMTrace initSt bound
```

Arbitrary OpRes: Prerequisites

Provided we know what state we are currently in, we can generate an operation and its result:

```
{currSt : ATMState} ->  
Arbitrary (resT : _ ** nsFn : resT -> ATMState **  
           OpRes resT currSt nsFn)  
where
```

Arbitrary OpRes: from Ready

From `Ready`, we can insert the card:

```
arbitrary {currSt=Ready} =  
  pure ( _ ** _ ** MkOpRes Insert () %search)
```

Arbitrary OpRes: from CardInserted

Using a dummy value for the PIN, we can control the frequencies of the getting the PIN right:

```
arbitrary {currSt=CardInserted} = do
  -- we need a PIN, even though we control the result
  let arbPIN = 0
  let op1 = (_ ** _ ** MkOpRes (CheckPIN arbPIN) Correct %search)
  let op2 = (_ ** _ ** MkOpRes (CheckPIN arbPIN) Incorrect %search)
  let op3 = (_ ** _ ** MkOpRes Eject () %search)

  frequency $ [(1, pure op1), (4, pure op2), (1, pure op3)]
```

Arbitrary OpRes: from Session

Generate an arbitrary amount to dispense, or eject the card:

```
arbitrary {currSt=Session} = do
  arbAmount <- arbitrary
  let op1 = (_ ** _ ** MkOpRes (Dispense arbAmount) ()) %search
  let op2 = (_ ** _ ** MkOpRes Eject ()) %search
  oneof [pure op1, pure op2]
```


Properties of the ATM

- Now that we have that, we can specify properties like “Within 5 state-transitions, we should be back in **Ready**”

```
public export 0
PROP_eventuallyReady : Fn (ATMTrace Ready 5) Bool
PROP_eventuallyReady =
  MkFn (\case (MkATMTrace _ trace) =>
        elem Ready (map (.resSt) trace))
```

- And* we can test it at compile-time

```
public export 0
EventuallyReady_OK : So (QuickCheck PROP_eventuallyReady)
EventuallyReady_OK = Oh
```

Model, verification, and implementation

- With most verification tools, we have to translate between models
 - Spec, model, and implementation are independent
- This facilitates translation mistakes
 - Might think we're verifying the same thing, when in actual fact the semantics have changed between representations

In our case, the specification *is* the model; *everywhere*

```
trace : (steps : Nat) -> (currSt : ATMState)
      -> Gen (Vect steps TraceStep)
trace 0 _ = pure []
trace (S k) currSt =
  the (Gen (x ** y ** OpRes x currSt y)) arbitrary
  >>=
  \case (_ ** nsFn ** opR@(MkOpRes _ res _)) =>
    do let nextState = nsFn res -- ! nsFn from ISM
       pure $ (MkTS opR nextState) :: !(trace k nextState)
```

QuickCheck spots the error!

- If we try to type-check the file we get:

```
-- Error: While processing right hand side of
--      EventuallyReady_OK. When unifying:
--      So True
-- and:
--      So (QuickCheck PROP_eventuallyReady)
-- Mismatch between: True and False.
```

- And if we investigate by running QC at the REPL, the error is exactly the fault in the model:

```
MkQCRes (Just False) <log> ""
Falsifiable, after 4 tests:
Starting @ Ready:
[ (<ATMOp 'Insert ~ ()'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
, (<ATMOp 'CheckPIN 0 ~ Incorrect'>, CardInserted)
]\n""
```

Fixing things

- Now that we know there's an error, we can fix things!

```
data ATMState
  = Ready
  | CardInserted Nat  -- track #retries
  | Session
  <...>
CheckPIN : (pin : Nat)
  -> ATMOp PINok (CardInserted (S tries))
      (\case Correct => Session
           Incorrect => ifThenElse (isZero tries)
                                   Ready
                                   (CardInserted tries))
```

- Carrying this through to the generators, our QC passes: file reloads successfully, the REPL reports

```
> QuickCheck PROP_eventuallyReady
MkQCRes (Just True) <log> "OK, passed 100 tests"
```

Benefits of a multifaceted approach

1. Adaptability — being able to use different tools
2. Speed — can trade speed for level of verification
 - This isn't about proving things, it is about increasing confidence in our typed models
3. **Coherence** — all done in one system
 - No need to translate to model-checking tool
 - Specification lives alongside model lives alongside implementation
 - The implementation is just there; it *is* runnable code
 - Parts can be verified independently *while* combined into an overall system

- As neat as this is, it is still convoluted to write the types, generators, etc.
- There are an abundance of FSM-like systems — the ARQ protocol, pick-and-place machines — which we plan to model
- This should hopefully reveal common patterns, which we can then factor out and automate large parts of this

Thank you



github.com/CodingCellist/talks

teh6@st-andrews.ac.uk